

# Claude Code Orchestrator

A Meta-Framework for AI-Driven Software Delivery

## EXECUTIVE OVERVIEW

**Prepared by:** Kearney

**Version:** 1.0

**Date:** November 2025

## Contents

- 1. Executive Summary**
- 2. Business Problems Solved**
- 3. Supported Project Types**
- 4. Key Features & Capabilities**
- 5. Technology Stack**
- 6. System Requirements**
- 7. Getting Started**
- 8. Target Users**

# Executive Summary

### KEY VALUE PROPOSITION

The Claude Code Orchestrator enables **structured, repeatable, high-quality project delivery at scale** - whether analytics pipelines, ML systems, web applications, or optimization models.

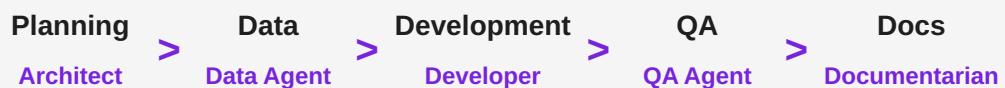
The **Claude Code Orchestrator** is a meta-framework that coordinates multiple specialized Claude Code subagents to collaboratively build complex software projects through a structured, checkpoint-driven workflow.

It solves the fundamental challenge of coordinating AI agents for enterprise-grade software delivery with quality assurance, governance compliance, and full traceability.



## Core Philosophy

Rather than using a single AI agent for all tasks, the orchestrator assigns specialized agents to phases where they excel:



Each phase concludes with validated checkpoint artifacts before proceeding.

# Business Problems Solved

---

## 1. Complexity Coordination

**Problem:** Complex projects require multiple phases with dependencies - architecture must be validated before coding, data must be processed before model training.

**Solution:** Checkpoint-driven workflow that sequences phases, validates completeness, and prevents downstream rework from bad early decisions.

## 2. Quality Assurance

**Problem:** Single agents can't consistently enforce quality standards across different project types and client requirements.

**Solution:** Automated quality gates (test coverage, security scanning, hygiene score) that must pass before phase progression.

## 3. Expertise Fragmentation

**Problem:** One AI agent can't excel at architecture design AND code implementation AND testing AND documentation.

**Solution:** Specialized subagents focused on their domain - Architect for design, Developer for code, QA for testing, Documentarian for guides.

## 4. Governance & Compliance

**Problem:** Client-specific compliance requirements (GDPR, HIPAA, SOC2, PCI-DSS) vary widely with no standardized enforcement.

**Solution:** Client governance system with YAML configurations that automatically enforce compliance requirements.

## 5. Decision Traceability

**Problem:** Complex projects involve many decisions; hard to answer "why did we choose this technology?"

**Solution:** Architecture Decision Records (ADRs) document context, rationale, and alternatives for significant choices.

## 6. Iteration & Rollback

**Problem:** Reverting errors after phase completion costs exponentially more effort.

**Solution:** Checkpoint artifacts enable non-destructive rollback to previous phases without losing progress.

# Supported Project Types

---

The orchestrator provides templates for four primary project types:

## Analytics Projects

**Best for:** Business intelligence, ETL pipelines, dashboards, reporting systems

<b>Tech Stack</b>	Python + SQL, pandas, DuckDB, Streamlit/Dash
-------------------	--

<b>Agents</b>	Architect, Data, Developer, QA, Documentarian
---------------	---

<b>Deliverables</b>	Processed datasets, dashboards, data quality reports
---------------------	--

*Example: Customer segmentation analysis, KPI dashboards, churn prediction reports*

## Machine Learning Projects

**Best for:** Model development, inference APIs, monitoring & retraining pipelines

<b>Tech Stack</b>	Python + scikit-learn/PyTorch, MLflow, FastAPI
-------------------	--

<b>Agents</b>	Architect, Data, Developer, QA, Documentarian
---------------	---

<b>Deliverables</b>	Trained models, inference APIs, monitoring dashboards
---------------------	---

*Example: Demand forecasting service, recommendation engine, fraud detection system*

## Web Applications

**Best for:** Full-stack apps, SPAs, admin dashboards, SaaS products

**Tech Stack**

TypeScript + React, FastAPI/Node.js, PostgreSQL

**Agents**

Architect, Developer, QA, Documentarian

**Deliverables**

Frontend code, backend APIs, tests, deployment guides

*Example: Customer portal, internal tool, executive dashboard*

## Supply Chain & Optimization

**Best for:** Operations research, resource allocation, scheduling, logistics

**Tech Stack**

Python + PuLP/Gurobi, pandas, scipy

**Agents**

Architect, Data, Optimization Engineer, Developer, QA, Documentarian

**Deliverables**

Optimization models, dashboards, scenario analysis

*Example: Route optimization, facility location, workforce scheduling*

# Key Features & Capabilities

---

## Multi-Agent Orchestration

Eight specialized agents with defined responsibilities:

Agent	Responsibility	Key Outputs
<b>Architect</b>	System design, technology selection	Architecture docs, ADRs
<b>Data</b>	ETL pipelines, feature engineering	Processed data, models, metrics
<b>Developer</b>	Feature implementation, APIs	Source code, tests
<b>QA</b>	Testing, validation	Test reports, coverage metrics
<b>Documentarian</b>	Documentation, guides	README, user guides, API docs
<b>Consensus</b>	Review and approve proposals	Approval decisions, feedback
<b>Steward</b>	Repository hygiene	Hygiene reports, cleanup
<b>Reviewer</b>	Code review (optional)	Review comments

## Client Governance

Automates enforcement of client-specific requirements:

### Quality Gates

Test coverage (50-95%),  
security scanning,  
complexity limits, peer  
review requirements

### Compliance

SOC2, GDPR, HIPAA,  
PCI-DSS, CCPA -  
automated enforcement

### Brand Constraints

Colors, fonts, terminology,  
logos - consistent  
branding

### Deployment Policies

Approval requirements,  
deployment windows,  
rollback policies

## Skills Framework

Reusable analytical methodologies:

- > **Time Series Analytics:** ARIMA, Prophet, LSTM for forecasting
- > **Optimization Modeling:** Linear programming, resource allocation
- > **Survey Data Processing:** Entity resolution, text analysis

Skills are domain-neutral - the same time series skill applies to telecom call forecasting, retail sales prediction, and healthcare admission planning.

## Architecture Decision Records (ADRs)

Document significant technical decisions with:

- Context and constraints
- Decision and rationale
- Alternatives considered

- Consequences and trade-offs

# Technology Stack

---

## Core Languages

Language	Version	Usage
Python	3.10+	Backend, data processing, ML
TypeScript	5.6+	Frontend, tooling
SQL	-	Data queries (DuckDB)

## Backend & API

- > **FastAPI** - Modern async web framework
- > **Uvicorn** - ASGI server
- > **Pydantic** - Data validation
- > **Typer** - CLI framework

## Data & Analytics

- > **pandas / polars** - Data manipulation
- > **DuckDB** - In-process OLAP database
- > **scikit-learn** - Machine learning
- > **matplotlib / plotly** - Visualization

## AI Integration

- > **Anthropic Claude API** - AI agent execution (claude-sonnet-4)

## Infrastructure & Deployment

- > **AWS Lambda** - Serverless compute
- > **AWS Amplify** - Frontend hosting
- > **GitHub Actions** - CI/CD automation
- > **Docker** - Containerization

## MLOps (Optional)

- > **MLflow** - Experiment tracking
- > **DVC** - Data version control
- > **Great Expectations** - Data quality

## Observability

- > **OpenTelemetry** - Distributed tracing
- > Custom metrics collection and dashboards

# System Requirements

---

## Development Environment

Component	Requirement
<b>Python</b>	3.10+ (3.11+ recommended)
<b>Node.js</b>	18.0+ (20.0+ for documentation site)
<b>npm</b>	9.0+
<b>Git</b>	2.0+
<b>Memory</b>	2GB+ for local development
<b>Disk</b>	5GB+ for development environment

## Required Configuration

```
# Environment variable for AI agent execution
export ANTHROPIC_API_KEY="your-api-key"

# Execution mode (optional, defaults to in_session)
export ORCHESTRATOR_EXECUTION_MODE="api"
```

## Installation

```
# Clone repository
git clone https://github.com/your-org/clause-code-orchestrator.git
cd clause-code-orchestrator

# Install Python dependencies
pip install -e .

# Install documentation site dependencies (optional)
cd site && npm install && cd ..
```

# Getting Started

---

## Step 1: Bootstrap a New Project

```
# Analytics project
orchestrator bootstrap analytics --output ~/projects/customer-analytics

# ML project
orchestrator bootstrap ml-model --output ~/projects/demand-forecast

# Web application
orchestrator bootstrap webapp --output ~/projects/admin-dashboard
```

## Step 2: Configure Project Intake

Edit the generated `intake.yaml`:

```
project:
  name: "Customer Analytics"
  type: "analytics"
goals:
  primary: ["Segment customers by value", "Identify churn risk"]
data:
  sources: ["Production database", "CRM API"]
  privacy_requirements: ["Anonymize PII"]
orchestration:
  enabled_agents: ["architect", "data", "developer", "qa", "documentarian"]
  consensus_required: ["planning", "data_engineering"]
```

## Step 3: Start the Workflow

```
# Start orchestration
orchestrator run start --intake intake.yaml

# Execute next phase
orchestrator run next

# Check status
orchestrator run status

# Approve consensus checkpoint
orchestrator run approve
```

## Step 4: Monitor Progress

```
# View metrics dashboard
orchestrator run metrics

# Run repository hygiene check
orchestrator run repo-hygiene
```

# Target Users

---

Persona	Use Case	Primary Benefit
<b>Data Scientists</b>	Analytics pipelines, ML models	Specialized Data agent, quality gates
<b>ML Engineers</b>	Production models, monitoring	Model versioning, drift detection
<b>Developers</b>	Full-stack applications	Architecture-first approach, test gates
<b>Consultants</b>	Client engagements	Governance compliance, ADR traceability
<b>OR Specialists</b>	Optimization problems	Dedicated optimization phase

## Best Fit Scenarios

- > **Complex multi-phase projects** requiring coordination across design, development, QA
  - > **Compliance-heavy engagements** (GDPR, HIPAA, SOC2 requirements)
  - > **Teams** needing clear roles and handoffs between phases
  - > **Repeatable workflows** across similar project types
  - > **Risk-sensitive projects** requiring quality gates and rollback capability
-

## Summary

The Claude Code Orchestrator transforms AI-assisted software development from ad-hoc single-agent interactions into a **structured, governance-compliant, checkpoint-driven workflow.**

### Specialization

Right agent for each phase

### Quality

Automated gates prevent mediocre deliverables

### Compliance

Client governance automatically enforced

### Traceability

Full audit trail via ADRs and checkpoints

#### KEY TAKEAWAY

Whether you're building a customer segmentation dashboard, deploying a demand forecasting API, or creating an executive reporting tool, the orchestrator provides the structure, quality gates, and documentation to deliver **production-grade results.**

Kearney Confidential

For questions, contact the platform team or open an issue in the repository.