

Contents

- 1. Executive Summary**
- 2. Business Problems Solved**
- 3. Supported Project Types**
- 4. Key Features & Capabilities**
- 5. Technology Stack**
- 6. System Requirements**
- 7. Getting Started**
- 8. Target Users**

Executive Summary

KEY VALUE PROPOSITION

The Claude Code Orchestrator enables **structured, repeatable, high-quality project delivery at scale** - whether analytics pipelines, ML systems, web applications, or optimization models.


The **Claude Code Orchestrator** is a meta-framework that coordinates multiple specialized Claude Code subagents to collaboratively build complex software projects through a structured, checkpoint-driven workflow.

It solves the fundamental challenge of coordinating AI agents for enterprise-grade software delivery with quality assurance, governance compliance, and full traceability.



Core Philosophy

Rather than using a single AI agent for all tasks, the orchestrator assigns specialized agents to phases where they excel:



Each phase concludes with validated checkpoint artifacts before proceeding.

Business Problems Solved

1. Complexity Coordination

Problem: Complex projects require multiple phases with dependencies - architecture must be validated before coding, data must be processed before model training.

Solution: Checkpoint-driven workflow that sequences phases, validates completeness, and prevents downstream rework from bad early decisions.

2. Quality Assurance

Problem: Single agents can't consistently enforce quality standards across different project types and client requirements.

Solution: Automated quality gates (test coverage, security scanning, hygiene score) that must pass before phase progression.

3. Expertise Fragmentation

Problem: One AI agent can't excel at architecture design AND code implementation AND testing AND documentation.

Solution: Specialized subagents focused on their domain - Architect for design, Developer for code, QA for testing, Documentarian for guides.

4. Governance & Compliance

Problem: Client-specific compliance requirements (GDPR, HIPAA, SOC2, PCI-DSS) vary widely with no standardized enforcement.

Solution: Client governance system with YAML configurations that automatically enforce compliance requirements.

5. Decision Traceability

Problem: Complex projects involve many decisions; hard to answer "why did we choose this technology?"

Solution: Architecture Decision Records (ADRs) document context, rationale, and alternatives for significant choices.

6. Iteration & Rollback

Problem: Reverting errors after phase completion costs exponentially more effort.

Solution: Checkpoint artifacts enable non-destructive rollback to previous phases without losing progress.

Supported Project Types

The orchestrator provides templates for four primary project types:

Analytics Projects

Best for: Business intelligence, ETL pipelines, dashboards, reporting systems

Tech Stack	Python + SQL, pandas, DuckDB, Streamlit/Dash
-------------------	--

Agents	Architect, Data, Developer, QA, Documentarian
---------------	---

Deliverables	Processed datasets, dashboards, data quality reports
---------------------	--

Example: Customer segmentation analysis, KPI dashboards, churn prediction reports

Machine Learning Projects

Best for: Model development, inference APIs, monitoring & retraining pipelines

Tech Stack	Python + scikit-learn/PyTorch, MLflow, FastAPI
-------------------	--

Agents	Architect, Data, Developer, QA, Documentarian
---------------	---

Deliverables	Trained models, inference APIs, monitoring dashboards
---------------------	---

Example: Demand forecasting service, recommendation engine, fraud detection system

Web Applications

Best for: Full-stack apps, SPAs, admin dashboards, SaaS products

Tech Stack

TypeScript + React, FastAPI/Node.js, PostgreSQL

Agents

Architect, Developer, QA, Documentarian

Deliverables

Frontend code, backend APIs, tests, deployment guides

Example: Customer portal, internal tool, executive dashboard

Supply Chain & Optimization

Best for: Operations research, resource allocation, scheduling, logistics

Tech Stack

Python + PuLP/Gurobi, pandas, scipy

Agents

Architect, Data, Optimization Engineer, Developer, QA, Documentarian

Deliverables

Optimization models, dashboards, scenario analysis

Example: Route optimization, facility location, workforce scheduling

Key Features & Capabilities

Multi-Agent Orchestration

Eight specialized agents with defined responsibilities:

Agent	Responsibility	Key Outputs
Architect	System design, technology selection	Architecture docs, ADRs
Data	ETL pipelines, feature engineering	Processed data, models, metrics
Developer	Feature implementation, APIs	Source code, tests
QA	Testing, validation	Test reports, coverage metrics
Documentarian	Documentation, guides	README, user guides, API docs
Consensus	Review and approve proposals	Approval decisions, feedback
Steward	Repository hygiene	Hygiene reports, cleanup
Reviewer	Code review (optional)	Review comments

Client Governance

Automates enforcement of client-specific requirements:

Quality Gates

Test coverage (50-95%),
security scanning,
complexity limits, peer
review requirements

Compliance

SOC2, GDPR, HIPAA,
PCI-DSS, CCPA -
automated enforcement

Brand Constraints

Colors, fonts, terminology,
logos - consistent
branding

Deployment Policies

Approval requirements,
deployment windows,
rollback policies

Skills Framework

Reusable analytical methodologies:

- > **Time Series Analytics:** ARIMA, Prophet, LSTM for forecasting
- > **Optimization Modeling:** Linear programming, resource allocation
- > **Survey Data Processing:** Entity resolution, text analysis

Skills are domain-neutral - the same time series skill applies to telecom call forecasting, retail sales prediction, and healthcare admission planning.

Architecture Decision Records (ADRs)

Document significant technical decisions with:

- Context and constraints
- Decision and rationale
- Alternatives considered

- Consequences and trade-offs

Technology Stack

Core Languages

Language	Version	Usage
Python	3.10+	Backend, data processing, ML
TypeScript	5.6+	Frontend, tooling
SQL	-	Data queries (DuckDB)

Backend & API

- > **FastAPI** - Modern async web framework
- > **Uvicorn** - ASGI server
- > **Pydantic** - Data validation
- > **Typer** - CLI framework

Data & Analytics

- > **pandas / polars** - Data manipulation
- > **DuckDB** - In-process OLAP database
- > **scikit-learn** - Machine learning
- > **matplotlib / plotly** - Visualization

AI Integration

- > **Anthropic Claude API** - AI agent execution (claude-sonnet-4)

Infrastructure & Deployment

- > **AWS Lambda** - Serverless compute
- > **AWS Amplify** - Frontend hosting
- > **GitHub Actions** - CI/CD automation
- > **Docker** - Containerization

MLOps (Optional)

- > **MLflow** - Experiment tracking
- > **DVC** - Data version control
- > **Great Expectations** - Data quality

Observability

- > **OpenTelemetry** - Distributed tracing
- > Custom metrics collection and dashboards

System Requirements

Development Environment

Component	Requirement
Python	3.10+ (3.11+ recommended)
Node.js	18.0+ (20.0+ for documentation site)
npm	9.0+
Git	2.0+
Memory	2GB+ for local development
Disk	5GB+ for development environment

Required Configuration

```
# Environment variable for AI agent execution
export ANTHROPIC_API_KEY="your-api-key"

# Execution mode (optional, defaults to in_session)
export ORCHESTRATOR_EXECUTION_MODE="api"
```

Installation

```
# Clone repository
git clone https://github.com/your-org/clause-code-orchestrator.git
cd clause-code-orchestrator

# Install Python dependencies
pip install -e .

# Install documentation site dependencies (optional)
cd site && npm install && cd ..
```

Getting Started

Step 1: Bootstrap a New Project

```
# Analytics project
orchestrator bootstrap analytics --output ~/projects/customer-analytics

# ML project
orchestrator bootstrap ml-model --output ~/projects/demand-forecast

# Web application
orchestrator bootstrap webapp --output ~/projects/admin-dashboard
```

Step 2: Configure Project Intake

Edit the generated `intake.yaml`:

```
project:
  name: "Customer Analytics"
  type: "analytics"
goals:
  primary: ["Segment customers by value", "Identify churn risk"]
data:
  sources: ["Production database", "CRM API"]
  privacy_requirements: ["Anonymize PII"]
orchestration:
  enabled_agents: ["architect", "data", "developer", "qa", "documentarian"]
  consensus_required: ["planning", "data_engineering"]
```

Step 3: Start the Workflow

```
# Start orchestration
orchestrator run start --intake intake.yaml

# Execute next phase
orchestrator run next

# Check status
orchestrator run status

# Approve consensus checkpoint
orchestrator run approve
```

Step 4: Monitor Progress

```
# View metrics dashboard
orchestrator run metrics

# Run repository hygiene check
orchestrator run repo-hygiene
```

Target Users

Persona	Use Case	Primary Benefit
Data Scientists	Analytics pipelines, ML models	Specialized Data agent, quality gates
ML Engineers	Production models, monitoring	Model versioning, drift detection
Developers	Full-stack applications	Architecture-first approach, test gates
Consultants	Client engagements	Governance compliance, ADR traceability
OR Specialists	Optimization problems	Dedicated optimization phase

Best Fit Scenarios

- > **Complex multi-phase projects** requiring coordination across design, development, QA
 - > **Compliance-heavy engagements** (GDPR, HIPAA, SOC2 requirements)
 - > **Teams** needing clear roles and handoffs between phases
 - > **Repeatable workflows** across similar project types
 - > **Risk-sensitive projects** requiring quality gates and rollback capability
-

Summary

The Claude Code Orchestrator transforms AI-assisted software development from ad-hoc single-agent interactions into a **structured, governance-compliant, checkpoint-driven workflow.**

Specialization

Right agent for each phase

Quality

Automated gates prevent mediocre deliverables

Compliance

Client governance automatically enforced

Traceability

Full audit trail via ADRs and checkpoints

KEY TAKEAWAY

Whether you're building a customer segmentation dashboard, deploying a demand forecasting API, or creating an executive reporting tool, the orchestrator provides the structure, quality gates, and documentation to deliver **production-grade results.**

Kearney Confidential

For questions, contact the platform team or open an issue in the repository.